

Downsampling Discrete Trajectories

Pascal Leroy (phl)

2026-01-26

This document describes the computations that are performed by the method `Append` of class `DiscreteTrajectorySegment` to downsample the points produced by an integrator and produce a compact yet accurate representation of discrete trajectories.

Overview

An integrator produces a stream of tuples (t_i, q_i, p_i) giving the degrees of freedom of a massless body at discrete times t_i . The purpose of downsampling is twofold:

- Construct cubic Hermite splines that interpolate between consecutive t_i to make it possible to evaluate the degrees of freedom at any time with sufficient accuracy.
- Make the representation more compact by dropping the tuples (t_i, q_i, p_i) for $i \in [1, n]$ if the Hermite spline constructed based on (t_1, q_1, p_1) and (t_n, q_n, p_n) approximates these tuples with sufficient accuracy.

A Brute Force Algorithm

The best way to describe the problem we want to solve is to present a brute-force algorithm that provides an exact solution:

Algorithm 1: BruteForceDownsampling.

Input: A tolerance η and a stream of tuples (t_i, q_i, p_i) .
Output: A stream of intervals $[j_1, j_2]$ such that the Hermite spline based on the tuples $(t_{j_1}, q_{j_1}, p_{j_1})$ and $(t_{j_2}, q_{j_2}, p_{j_2})$ approximates the input tuples (t_i, q_i, p_i) for $i \in [j_1, j_2]$ with a tolerance better than η .

1. Let $j \leftarrow 1$.
2. **while** *not at end of input stream* **do**
3. Construct the Hermite spline h based on the tuples (t_j, q_j, p_j) and (t_i, q_i, p_i) .
4. **if** $\max_{k=j \dots i} \|h(t_k) - q_k\|_2 > \eta$ **then**
5. Emit the interval $[j, i - 1]$.
6. Let $j \leftarrow i - 1$.
7. **end**
8. Emit the interval $[j, n]$ where n is the index of the last element in the stream.

A few things are worth noting here:

- Only the tuples $(t_{j_1}, q_{j_1}, p_{j_1})$ and $(t_{j_2}, q_{j_2}, p_{j_2})$ corresponding to the bounds of the intervals $[j_1, j_2]$ need to be stored permanently. The other tuples can be interpolated with an accuracy better than η .
- The Hermite splines do not need to be stored permanently; they can be reconstructed based on the tuples $(t_{j_1}, q_{j_1}, p_{j_1})$ and $(t_{j_2}, q_{j_2}, p_{j_2})$.
- The algorithm is optimal in the sense that it produces the longest possible intervals $[j_1, j_2]$ that satisfy the tolerance η .

- For each input tuple, the algorithm needs to scan all past tuples since the upper bound j of the last emitted interval. Therefore, the algorithm is quadratic in the number of tuples in the input stream.

A Binary Search Algorithm

In order to avoid the quadratic complexity, we used to use a binary search algorithm as follows:

Algorithm 2: BinarySearchDownsampling.

Input: A tolerance η , an integer N , and a stream of tuples (t_i, q_i, p_i) .
Output: A stream of intervals $[j_1, j_2]$ such that the Hermite spline based on the tuples $(t_{j_1}, q_{j_1}, p_{j_1})$ and $(t_{j_2}, q_{j_2}, p_{j_2})$ approximates the input tuples (t_i, q_i, p_i) for $i \in [j_1, j_2]$ with a tolerance better than η .

```

1. Let  $A$  be an array of size  $N$  used to store tuples.
2. while not at end of input stream do
3.   if  $A$  is full then
4.     Let  $k \leftarrow \lfloor N/2 \rfloor$ .
5.     Emit SearchInterval( $A, 1, k$ ).
6.     Emit SearchInterval( $A, k + 1, N$ ).
7.     Clear  $A$ .
8.   Append  $(t_i, q_i, p_i)$  to  $A$ .
9. end
10. Function SearchInterval( $A, i_1, i_2$ ) is
11.   Construct the Hermite spline  $h$  based on the tuples  $(t_{i_1}, q_{i_1}, p_{i_1})$  and  $(t_{i_2}, q_{i_2}, p_{i_2})$ .
12.   if  $\max_{k=i_1 \dots i_2} \|h(t_k) - q_k\|_2 > \eta$  then
13.     Let  $k \leftarrow \lfloor i_2/2 \rfloor$ .
14.     Emit SearchInterval( $A, i_1, k$ ).
15.     Emit SearchInterval( $A, k + 1, i_2$ ).
16.   else
17.     Emit  $[i_1, i_2]$ .
18.   end
19. end
```

This algorithm has the following properties:

- The array A must be stored permanently since it contains tuples that have not been downsampled yet and that will be needed to make a future decision about the intervals to emit.
- The choice of N involves a trade-off: N must be large enough that the down-sampling is effective and produces long intervals $[j_1, j_2]$; but it must not be so large that the size of A affects performance.
- The algorithm has a complexity of $\mathcal{O}(N \log N)$ on average and $\mathcal{O}(N^2)$ in the worst case.
- The algorithm does not produce the longest possible intervals: in the worst case, all the emitted intervals may be too small by a factor of 2.

An Efficient Streaming Algorithm

In this section we present an algorithm that, like the brute-force algorithm, is streaming (so a decision is made on each input tuple without having to resort to intermediate storage), whose complexity is linear in the size of the stream, and which generally produces a much more compact output than the binary search algorithm. Without loss of generality we assume that, when the tuple of index i is received, no interval was ever emitted (so the first interval to emit will necessarily have the lower bound 1).

The algorithm is incremental, so let's assume that we have already constructed a Hermite spline h_{i-1} with an error bound ε_{i-1} such that:

$$\max_{k=1 \dots i-1} \|h_{i-1}(t_k) - q_k\|_2 \leq \varepsilon_{i-1} \leq \eta$$

It is obviously possible to build such a spline, because the error on a spline build based on two consecutive points in the stream is 0. In other words, $\varepsilon_2 = 0$ (ε_1 is not defined because a spline cannot be constructed using a single point).

When the point with index i is received, we construct the Hermite spline h_i based on the tuples (t_1, q_1, p_1) and (t_i, q_i, p_i) . Noting that $h_i(t_i) = q_i$ we have:

$$\max_{k=1 \dots i} \|h_i(t_k) - q_k\|_2 = \max_{k=1 \dots i-1} \|h_i(t_k) - q_k\|_2$$

And, by the triangular inequality:

$$\max_{k=1 \dots i-1} \|h_i(t_k) - q_k\|_2 \leq \max_{k=1 \dots i-1} \|h_i(t_k) - h_{i-1}(t_k)\|_2 + \max_{k=1 \dots i-1} \|h_{i-1}(t_k) - q_k\|_2$$

The first term is bounded by the L^∞ norm of $h_i - h_{i-1}$ and the second term is bounded, because of our recurrence hypothesis, by ε_{i-1} . Thus:

$$\max_{k=1 \dots i} \|h_i(t_k) - q_k\|_2 \leq \|h_i - h_{i-1}\|_\infty + \varepsilon_{i-1}$$

Therefore, if we can compute $\|h_i - h_{i-1}\|_\infty$ we can find an upper bound ε_i on the left-hand of this inequality and use it to determine if h_i is still within the tolerance η or if the interval must be split.

Unfortunately $h_{i-1}(t)$ and $h_i(t)$ are 3rd degree polynomials, so $\|h_i(t) - h_{i-1}(t)\|_2^2$ is a 6th degree polynomial, and finding its extrema requires the resolution of a 5th degree equation, which would probably be overly expensive.

We can note however that, by construction, the h_i are exact for $t = t_i$. Therefore, $h_{i-1}(t_1) = h_i(t_1)$ and $h'_{i-1}(t_1) = h'_i(t_1)$. This implies that $(t - t_1)^2$ divides the polynomial $h_i(t) - h_{i-1}(t)$. There exist a 1st degree polynomial $q_i(t)$ such that:

$$h_i(t) - h_{i-1}(t) = (t - t_1)^2 q_i(t)$$

Switching to norms we find:

$$\|h_i(t) - h_{i-1}(t)\|_2^2 = (t - t_1)^4 \|q_i(t)\|_2^2$$

To find the extrema we need to compute the roots of the derivative:

$$\begin{aligned} \frac{d}{dt} \|h_i(t) - h_{i-1}(t)\|_2^2 &= 4(t - t_1)^3 \|q_i(t)\|_2^2 + (t - t_1)^4 (q_i(t) \cdot q'_i(t)) \\ &= 2(t - t_1)^3 (q_i(t) \cdot (2q_i(t) + (t - t_1)q'_i(t))) \end{aligned}$$

Now the polynomial $q_i(t) \cdot (2q_i(t) + (t - t_1)q'_i(t))$ is a 2nd degree polynomial and its roots may be computed efficiently. In order to compute the extrema of $h_i(t) - h_{i-1}(t)$, we must evaluate it at the roots of its derivative, and at t_{i-1} , since it could have an extremum there that does not correspond to a zero of the derivative.