

Downsampling Discrete Trajectories

Pascal Leroy (phl)

2026-01-26

This document describes the computations that are performed by the method `Append` of class `DiscreteTrajectorySegment` to downsample the points produced by an integrator and produce a compact yet accurate representation of discrete trajectories.

Overview

An integrator produces a stream of tuples (t_i, q_i, p_i) giving the degrees of freedom of a massless body at discrete times t_i . The purpose of downsampling is twofold:

- Construct cubic Hermite splines that interpolate between consecutive t_i to make it possible to evaluate the degrees of freedom at any time with sufficient accuracy.
- Make the representation more compact by dropping the tuples (t_i, q_i, p_i) for $i \in [1, n]$ if the Hermite spline constructed based on (t_1, q_1, p_1) and (t_n, q_n, p_n) approximates these tuples with sufficient accuracy.

A Brute Force Algorithm

The best way to describe the problem we want to solve is to present a brute-force algorithm that provides an exact solution:

Algorithm 1: BruteForceDownsampling.

Input: A tolerance η and a stream of tuples (t_i, q_i, p_i) .
Output: A stream of intervals $[j_1, j_2]$ such that the Hermite spline based on the tuples $(t_{j_1}, q_{j_1}, p_{j_1})$ and $(t_{j_2}, q_{j_2}, p_{j_2})$ approximates the input tuples (t_i, q_i, p_i) for $i \in [j_1, j_2]$ with a tolerance better than η .

1. Let $j = 1$.
 2. **while** *not at end of input stream* **do**
 3. | Compute the Hermite spline h based on the tuples (t_j, q_j, p_j) and (t_i, q_i, p_i) .
 4. | **if** $\sum_{k=j}^i \|h(t_k) - q_k\|_2 > \eta$ **then**
 5. | | Emit the interval $[j, i - 1]$.
 6. | | Let $j = i - 1$.
 7. **end**
 8. Emit the interval $[j, n]$ where n is the index of the last element in the stream.
-

A few things are worth noting here:

- Only the tuples $(t_{j_1}, q_{j_1}, p_{j_1})$ and $(t_{j_2}, q_{j_2}, p_{j_2})$ corresponding to the bounds of the intervals $[j_1, j_2]$ need to be stored permanently. The other tuples can be interpolated with an accuracy better than η .
- The Hermite splines do not need to be stored permanently; they can be reconstructed based on the tuples $(t_{j_1}, q_{j_1}, p_{j_1})$ and $(t_{j_2}, q_{j_2}, p_{j_2})$.
- The algorithm is optimal in the sense that it produces the longest possible intervals $[j_1, j_2]$ that satisfy the tolerance η .

- For each input tuple, the algorithm needs to scan all past tuples since the upper bound j of the last emitted interval. Therefore, the algorithm is quadratic in the number of tuples in the input stream.